

Question 3 + 4

720017170

Question 3 - See QMD for Code

(a) Import each dataset into memory as a separate data frame, keeping all countries as your sample.

```
plt.rcParams.update({'font.size': 14})

# Loading in the Data
Health_Data = pd.read_csv('Data/Health.csv', index_col=None)
Infant_Data = pd.read_csv("Data/Infant.csv", index_col=None)

# Replace .. with NA
Health_Data.replace("..", pd.NA, inplace=True)
Infant_Data.replace("..", pd.NA, inplace=True)

# Removing unnecessary columns
Health_Data = Health_Data.drop(columns=['Series Name', 'Series Code'])
Infant_Data = Infant_Data.drop(columns=['Series Name', 'Series Code'])

# Remove names in []
Health_Data.columns = Health_Data.columns.str.replace(r'\[.*\]', '', regex=True)
Infant_Data.columns = Infant_Data.columns.str.replace(r'\[.*\]', '', regex=True)

print(Infant_Data.head())
print(Health_Data.head())
```

	Country Name	Country Code	2000	2001	2002	2003	2004	2005	2006	\
0	Afghanistan	AFG	92	89.3	86.6	83.7	80.9	78	75.1	
1	Albania	ALB	24	22.9	21.6	20.4	19.1	17.8	16.5	
2	Algeria	DZA	35.6	34.3	33	31.6	30.3	29	27.8	
3	American Samoa	ASM	<NA>	<NA>	<NA>	<NA>	<NA>	<NA>	<NA>	

4	Andorra	AND	6.5	6.3	6	5.8	5.6	5.3	5.1
---	---------	-----	-----	-----	---	-----	-----	-----	-----

	2007	...	2014	2015	2016	2017	2018	2019	2020	2021	2022	2023
0	72.3	...	56.2	54.6	53	51.5	50.1	48.8	47.4	46.1	44.8	<NA>
1	15.3	...	8.8	8.5	8.4	8.3	8.3	8.3	8.4	8.4	8.4	<NA>
2	26.6	...	22	21.7	21.4	21	20.6	20.1	19.7	19.2	18.7	<NA>
3	<NA>	...	<NA>	<NA>	<NA>	<NA>	<NA>	<NA>	<NA>	<NA>	<NA>	<NA>
4	4.9	...	3.5	3.4	3.2	3.1	3	2.9	2.8	2.7	2.6	<NA>

[5 rows x 26 columns]

	Country Name	Country Code	2000	2001	2002	\
0	Afghanistan	AFG	<NA>	<NA>	17.00758553	
1	Albania	ALB	65.1501236	73.78884125	78.99478149	
2	Algeria	DZA	62.11769485	67.33850098	66.94760132	
3	American Samoa	ASM	<NA>	<NA>	<NA>	
4	Andorra	AND	1287.00280762	1336.21142578	1486.171875	

	2003	2004	2005	2006	2007	\
0	17.81492424	21.42946434	25.10707283	28.91982269	32.71720505	
1	106.29218292	138.11340332	152.12762451	166.81382751	212.61096191	
2	76.23547363	93.02433014	101.30373383	117.43313599	151.77920532	
3	<NA>	<NA>	<NA>	<NA>	<NA>	
4	1772.71337891	1990.0748291	2214.64697266	2139.27539063	2489.43115234	

	...	2014	2015	2016	2017	\
0	...	60.18957901	60.05854034	61.48645782	66.90921783	
1	...	295.12359619	255.35635376	277.04321289	297.4619751	
2	...	361.15942383	292.275177	261.40023804	265.83843994	
3	...	<NA>	<NA>	<NA>	<NA>	
4	...	3089.84301758	2688.20629883	2755.44848633	2873.29614258	

	2018	2019	2020	2021	2022	2023
0	71.33430481	74.23410797	80.28805542	81.31976318	<NA>	<NA>
1	351.3012085	367.75839233	396.88024902	464.74285889	<NA>	<NA>
2	266.46469116	235.99041748	206.03512573	204.56661987	<NA>	<NA>
3	<NA>	<NA>	<NA>	<NA>	<NA>	<NA>
4	3164.38842773	3026.59741211	3269.29736328	3505.99145508	<NA>	<NA>

[5 rows x 26 columns]

(b) If data are not already stored in this way, please reshape data so that they consist of a single line of data for each country and year.

```
# Pivoting the Data into a long format
Health_Data_long = pd.melt(Health_Data,
    id_vars=['Country Name', 'Country Code'],
    var_name='Year',
    value_name='Heathcare Expenditure (USD)')
Infant_Data_long = pd.melt(Infant_Data,
    id_vars=['Country Name', 'Country Code'],
    var_name='Year',
    value_name='Infant Mortality Rates (per 1,000 live births)')
```

	Country Name	Country Code	Year	Heathcare Expenditure (USD)
0	Afghanistan	AFG	2000	<NA>
1	Albania	ALB	2000	65.1501236
2	Algeria	DZA	2000	62.11769485
3	American Samoa	ASM	2000	<NA>
4	Andorra	AND	2000	1287.00280762

	Country Name	Country Code	Year	Infant Mortality Rates (per 1,000 live births)
0	Afghanistan	AFG	2000	92
1	Albania	ALB	2000	24
2	Algeria	DZA	2000	35.6
3	American Samoa	ASM	2000	<NA>
4	Andorra	AND	2000	6.5

(c) Calculate the total number of countries observed in each data frame Calculate the total number of years observed in each data frame.

```
# Counts the number of unique contries
num_countries_FDI = Health_Data_long["Country Name"].nunique()

# Outputs the number of unique countries using an f string
print(f"Total number of unique countries in Health_Data: {num_countries_FDI}")

num_years_FDI = Health_Data_long["Year"].nunique() # Counts the number of unique years
print(f"Total number of unique years observed in Health_Data: {num_years_FDI}")
```

```

num_countries_GDP = Infant_Data_long["Country Name"].nunique()
print(f"Total number of unique countries in Infant_Data: {num_countries_GDP}")

num_years_GDP = Infant_Data_long["Year"].nunique()
print(f"Total number of unique years observed in Infant_Data: {num_years_GDP}")

```

```

Total number of unique countries in Health_Data: 217
Total number of unique years observed in Health_Data: 24
Total number of unique countries in Infant_Data: 217
Total number of unique years observed in Infant_Data: 24

```

(d) Calculate the number of observations for which data is missing

```

# Sums the number of missing values in each dataset
missing_values_Health = Health_Data_long.isna().sum().sum()
print(f"Total missing observations in Health_Data: {missing_values_Health}")

missing_values_Infant = Infant_Data_long.isna().sum().sum()
print(f"Total missing observations in Infant_Data: {missing_values_Infant}")

```

```

Total missing observations in Health_Data: 1099
Total missing observations in Infant_Data: 700

```

(e) Join the two files by country and year so that you have single dataframe containing both variables. Explain clearly what type of join this is, and carefully check that the number of observations resulting from the join makes sense.

```

# Merge the data on Country Name, Country code and Year
merged_data = pd.merge(Health_Data_long, Infant_Data_long, on=['Country Name',
'Country Code', 'Year'])
print(merged_data.head())

# Print the number of rows in the DataFrame
num_rows = merged_data.shape[0]
print(f"Number of rows in the DataFrame: {num_rows}")

```

	Country Name	Country Code	Year	Healthcare Expenditure (USD)	\
0	Afghanistan	AFG	2000	<NA>	
1	Albania	ALB	2000	65.1501236	
2	Algeria	DZA	2000	62.11769485	

3	American Samoa	ASM	2000	<NA>
4	Andorra	AND	2000	1287.00280762

Infant Mortality Rates (per 1,000 live births)	
0	92
1	24
2	35.6
3	<NA>
4	6.5

Number of rows in the DataFrame: 5208

The join completed in the above code chunk is an inner join and only keeps rows that exist in both `Health_Data_long` and `Infant_Data_long`. If a country-year exists in one dataset but not the other, it will be dropped.

Question 4 - Investigating the Relationship Between Current Healthcare Expenditure per Capita and Infant Mortality Rates from 2000 - 2022

Missing Data

Table 1: Missing Observations of Variables

Dataset	Total Missing Observations	Percentage of Missing Observations
Health Expenditure Dataset	1099	21.1
Infant Mortality Rate Dataset	700	13.4

Both datasets contained a large amount of missing data, illustrated in Table 1. Potentially due to countries not collecting the data or collecting the data at different year intervals. Missing data can have a large impact on data analysis if not handled properly and can lead to skewed or incorrect conclusions. The year 2023 contained no data; therefore, this column was dropped. To deal with the other missing data, I decided to drop all rows containing missing data, sometimes, this could result in a significant reduction of sample size; however, in this case with observational data, 1099 observations were removed (21.1% of the dataset) and only 27 countries were dropped, indicating this was an effective method to handling missing data as there were still 4109 observations. An alternative approach would've been mean, multiple or regression imputation if dropping rows with missing data caused a significant decrease in sample size.

Summary Statistics

Table 2: Summary Statistics of Variables

Variable	N	Mean	Median	SD	Min	Max
Healthcare Expenditure (USD)	4109.0	956.0	256.7	1685.7	4.0	12473.8
Infant Mortality Rates (per 1,000 live births)	4109.0	26.9	17.7	25.0	1.4	138.3

Table 2 displays the summary statistics for healthcare expenditure (USD) and infant mortality rates (per 1,000 live births) across 4109 observations, revealing significant differences between countries.

Healthcare expenditure per Capita showed a mean of \$956.0 but a lower median of \$256.7, indicating a negatively skewed distribution where only few countries spend more. The large standard deviation (\$1,685.7) and range (\$4.0–\$12,473.8) highlight large global and temporal differences in healthcare investment.

Infant mortality rates show similar variation, with a mean of 26.9 deaths per 1,000 live births and a median of 17.7. The high standard deviation (25.0) and range (1.4–138.3) may be attributed to major differences in healthcare investment and quality.

Distribution Analysis




Figure 1: Box Plot of Healthcare expenditure Per Capita (USD)




Figure 2: Histogram with Density of Healthcare expenditure Per Capita (USD)

Figures 1 and 2 show the distribution of healthcare expenditure per capita (USD). Figure 1 shows that healthcare expenditure is highly negatively skewed, supporting the analysis from the summary statistics. The median expenditure is toward the lower quartile, indicating that most countries have relatively little expenditure, while a few have significantly higher spending. The whiskers of the box plot are short, suggesting that a large proportion of the data is concentrated within a lower range, while the numerous outliers highlight extreme expenditure levels in some countries.

Figure 2 reiterates the negative skew of the data. Most countries have low healthcare expenditure, grouped toward the left of the axis, with just a handful having exceptionally high expenditures. The density curve (smooth blue line) shows the exponential drop in frequency as expenditure increases, emphasising that high-spending countries are exceptions rather than the rule.




Figure 3: Box Plot of Infant Mortality Rate (per 1,000 live births)




Figure 4: Histogram with Density of Infant Mortality Rate (per 1,000 live births)

The distribution of the infant mortality rate (per 1,000 live births) seen in Figures 3 and 4 is similar to that of healthcare expenditure per capita (USD).

There are several outliers in Figure 3, with specific countries having abnormally high infant mortality rates. The variable is again negatively skewed, as seen by the median being significantly lower than the upper quartile.

This is supported by Figure 4, which displays a dramatic fall in frequency as rates rise, with most values clustering below 40 deaths per 1,000 live births. While infant mortality is low in many nations, it is much higher in others, most likely because of infrastructural constraints, economic considerations, and healthcare discrepancies.

Due to the number of outliers within both variables, I decided to use the Interquartile Range (IQR) method to remove any outliers to ensure the regression analysis is more reliable and accurate.

Outliers may skew the models fit and influence regression coefficients, inflating evaluation metrics such as mean absolute error (MAE). By applying the IQR method, the data represents a more consistent trend, reducing the impact of extreme values and improving the model's ability to capture the true relationship between healthcare expenditure and infant mortality rates.

Correlation Analysis




Figure 5: Spearman Rank Correlation Matrix of Healthcare Expenditure and Infant Mortality Rates

Figure 5 illustrates the relationship between healthcare expenditure and infant mortality rates. The correlation coefficient of -0.82 indicates a strong negative correlation, suggesting that as healthcare expenditure increases, infant mortality rates to decrease. This aligns with economic and public health expectations, where greater investment in healthcare typically leads to better medical infrastructure, improved care, and reduced infant deaths. Spearman’s rank correlation is used as it captures non-linear relationships, making it more robust. However, correlation does not imply causation, and additional factors such as healthcare efficiency, socioeconomic disparities, and government policies could be confounders in this relationship.

Regression Analysis




Figure 6: Comparison of Regression Models: Healthcare Spend vs Infant Mortality Rate

Figure 6 compares five regression models; exponential, polynomial, logarithmic, power law, and linear, in quantifying the relationship between healthcare expenditure per capita and infant mortality rates. All models are constrained to prevent infant mortality predictions below 0, ensuring a realistic representation. The linear model performs worst ($R^2 = 0.33$), not capturing the non-linearity of the data, indicating that non-linear regression methods may fit the relationship better.

Polynomial regression initially follows the relationship but overall performs poorly, shown by its lower R^2 of 0.48. Exponential and logarithmic regression perform well ($R^2 = 0.56$ and 0.61, respectively), capturing the steep initial decline in infant mortality rates before plateauing at higher expenditure levels. The power law model achieves the best fit with an R^2 of 0.61,

revealing that 61% of the variance in infant mortality rates is due to healthcare expenditure per capita, and the lowest mean absolute error (MAE) of 10.47.

This analysis proves that non-linear models, particularly power law and logarithmic regression, provide the most accurate representation of the data, reiterating the non-linearity of the relationship between healthcare expenditure per capita and infant mortality rates shown by Figure 5.

Regression Model Evaluation




Figure 7: Residual Plot of Power Law Regression

The residuals, illustrated by Figure 7, for the power law regression model indicates some issues with fit of the model. The residuals display a clear pattern rather than being randomly scattered around 0, suggesting heteroscedasticity, where the variance of residuals increases as fitted values increase. This shows that the model performs well at low levels of healthcare expenditure but struggles to maintain accuracy as expenditure increases. To address these issues, applying a log transformation to the dependent variable or using alternative regression techniques may help improve the model's fit.

Time Series Analysis

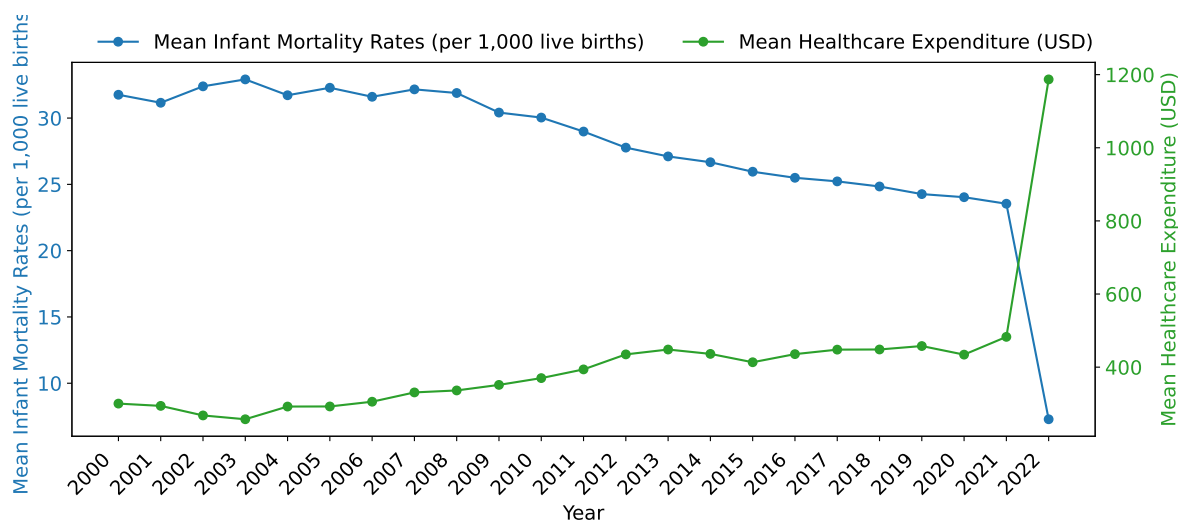


Figure 8: Time Series Analysis of Mean Infant Mortality Rates and Healthcare Expenditure

Figure 8 illustrates the non-linear negative association between mean infant mortality rates (IMR) and mean healthcare expenditure (USD) from 2000 to 2022. The trend indicates healthcare expenditure has steadily climbed, indicating greater investment in health infrastructure and services. Leading to a continual drop in IMR, reflecting these improvements in healthcare, economic development, and medical advances

An anomaly occurs in 2022 when healthcare expenditure increases disproportionately to other years and infant death rates significantly decrease. This large increase is likely due to pandemic-related expenditure, emergency health interventions, or data errors.

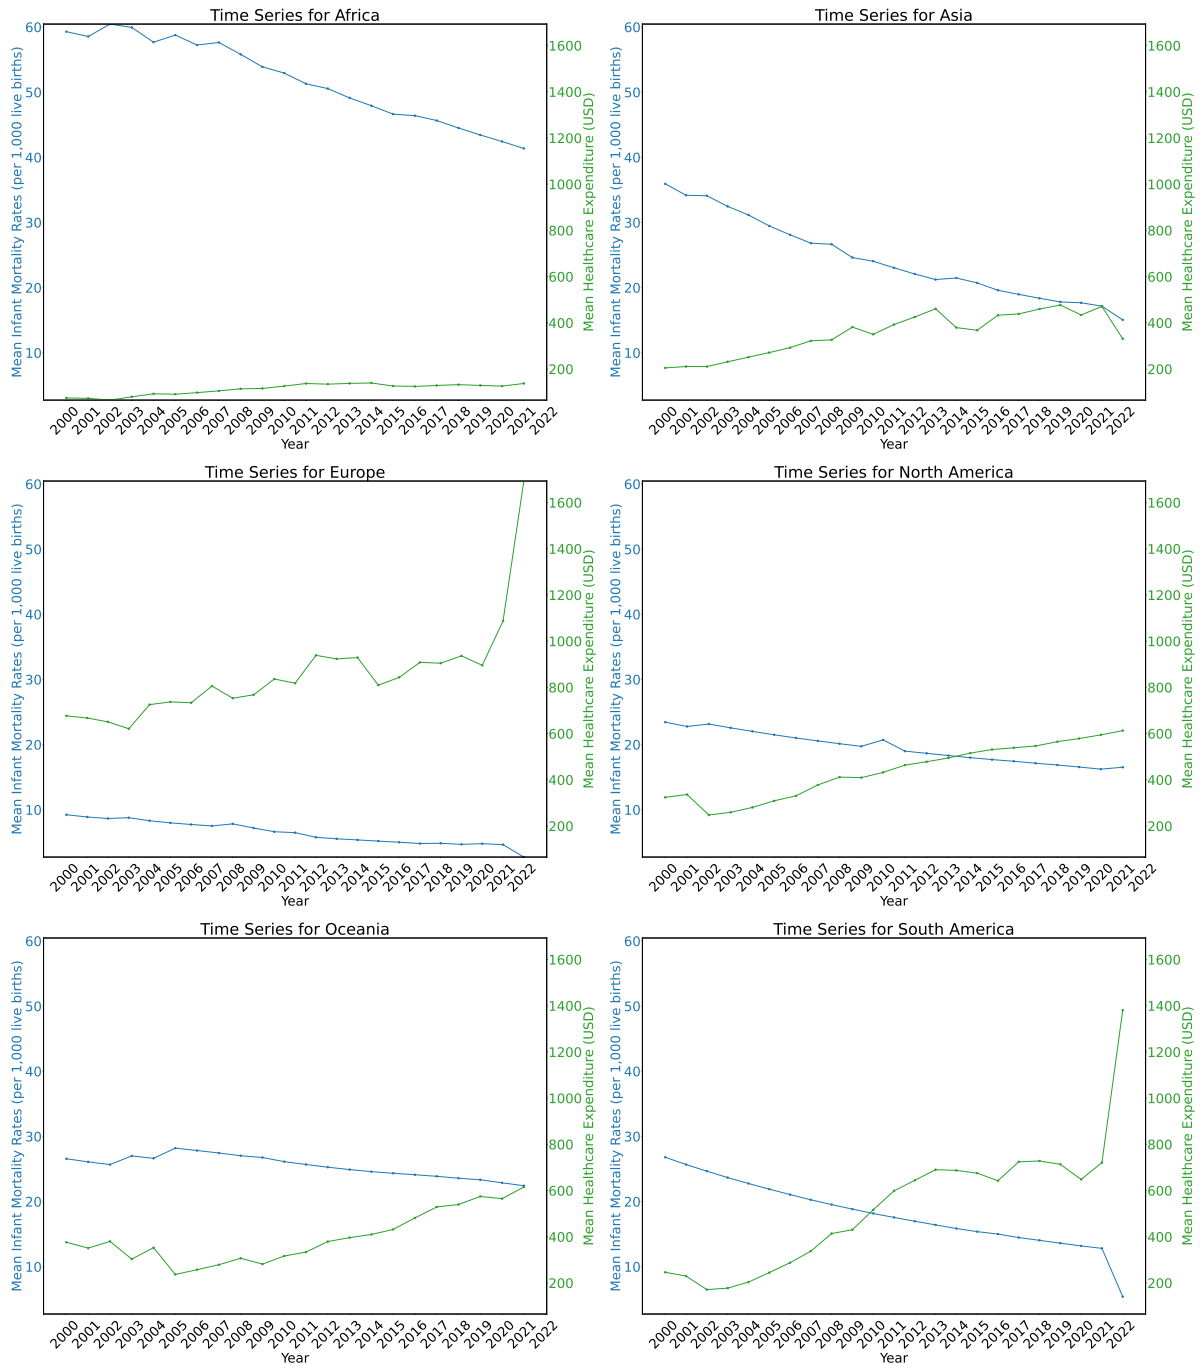


Figure 9: Time Series Analysis of Mean Infant Mortality Rates and Healthcare Expenditure By Continent

Figure 9 shows continent-specific relationships between IMR and mean healthcare expenditure (USD) from 2000 to 2022. The constant negative relationship across all continents reiterates that increasing healthcare investment has the potential to reduce infant mortality.

Asia and Africa have had considerable decreases in IMR, indicating substantial improvements in healthcare despite relatively small expenditure. Europe and North America, with larger starting expenditures, experienced lower IMR decreases, indicating diminishing returns on investment.

Healthcare expenditure rose significantly in 2022 in Europe and South America, potentially due to pandemic-related measures. This substantial rise in expenditure corresponds with an acceleration in IMR decreases, implying short-term healthcare gains.

Conclusion

In conclusion, there is a strong non-linear negative relationship between healthcare expenditure and infant mortality rates across global regions. Countries with higher healthcare investment generally experience lower infant mortality, though the impact varies based on economic and healthcare infrastructure factors.

The non-linear nature of this relationship indicates diminishing returns, where initial increases in spending lead to substantial improvements, but further investments yield smaller reductions in infant mortality. Regression and time-series analyses further reinforce this pattern, indicating long-term declines in infant mortality alongside growing healthcare expenditure.

Regional disparities are present, with high-income regions investing more per capita while achieving lower mortality rates, whereas lower-income regions show greater relative improvements despite lower absolute expenditure. The 2022 anomaly suggests short-term shifts in healthcare spending and outcomes, likely due to pandemic-driven policies.

While healthcare expenditure is an important driver of infant mortality rates, efficient allocation, accessibility, and policy effectiveness remain key determinants of long-term health improvements worldwide.

[Link to Github Repository = BEE2041 Data Science in Economics Assignment](#)